

Agentic Ops Benchmark

Local 30B-class models on a single consumer GPU, measured against hosted frontier models on the same agentic coding and operations tasks — under two different agent harnesses.

Independent single-operator evaluation · 2026-08-12 · 7 models × 15 tasks × 2 harnesses = 210 graded runs

96.7%

BEST LOCAL PASS RATE

0.844

BEST LOCAL QUALITY

51s

FASTEST LOCAL
MEDIAN

1 day

FROM RELEASE TO
BENCHMARKED

■ local (self-hosted) ■ frontier (hosted API)

The stack

Everything local runs on one node: a Proxmox host (H15) with an **NVIDIA RTX 3090, 24 GB**, attached as an external GPU over Thunderbolt 4. Models are served with **llama.cpp** and hot-swapped on demand — only one is resident at a time — behind a **LiteLLM** gateway that presents a single OpenAI-compatible endpoint to every agent. Frontier models are reached over their vendor APIs. Both agent harnesses are pointed at the same gateway, so a model is addressed identically no matter which harness is driving it.

Two of the four local models were benchmarked within days of their release: **Muse Glimmer 30B** (Meta Superintelligence Labs, 10 Aug 2026) and **Nemotron 3.5 Lightning 30B-A3B** (NVIDIA, 11 Aug 2026) — the latter one day after it shipped. **Prime Agent** (Prime Intellect, open-sourced 6 Aug 2026) was six days old.

How this was measured

Every model runs the same 15 tasks twice — once driven by each of two agent harnesses — inside an isolated working copy, with no access to the grading tests as an oracle. The agent is told to run `pytest` and iterate until green. Each run is scored on two independent axes:

Pass rate is the mean per-test pass fraction of the hidden suite, so partial credit is real: fixing one of two planted bugs scores 0.5. **Quality** is a 0–1 LLM-as-judge score rating the finished solution for correctness, edge-case handling and craft — a fragile pass scores below a clean one. **Latency** is wall-clock time to a finished task, including every tool call and retry the agent made, not token throughput.


On the judge. The judge is GPT-5.5, which is also one of the models under test. A model cannot credibly grade its own work, so **GPT-5.5's own quality scores are withheld** and shown as *n/a* throughout; its pass rate, which is decided by `pytest` and not by any model's opinion, is reported normally. Every other model's quality score comes from a system that is not itself competing in that row.

The task battery was curated with Fable 5 and mirrors my day-to-day agentic work: AIOps alert correlation, NetDevOps intent-to-ACL compilation and config-drift planning, HPC backfill scheduling, inference KV-budget batching, RAG pipeline debugging and tool-calling loops — alongside held-out EvalPlus coding problems, cross-module refactors and long-context needle retrieval.

Leaderboard


Quality index

Mean LLM-as-judge score across all 30 runs per model. Higher is better. GPT-5.5 is absent because it is the judge.



Pass rate

Mean per-test pass fraction of the hidden suites, with partial credit.




#	MODEL	CLASS	QUALITY	PASS RATE	MEDIAN LATENCY	SERVED AS
1	GPT-5.6 Sol OpenAI, GA 9 Jul 2026	FRONTIER	0.913	100.0%	40s	hosted API
2	Claude Opus 4.8 Anthropic	FRONTIER	0.907	100.0%	56s	hosted API
3	GPT-5.5 OpenAI	FRONTIER	n/a	98.0%	52s	hosted API
4	Qwen3.6 27B Dense 27B, MTP draft head	LOCAL	0.844	96.7%	64s	Q4_K_M · 16 GB · 98,304 ctx
5	Muse Glimmer 30B Meta Superintelligence Labs, 30B · launched 10 Aug 2026	LOCAL	0.844	96.7%	272s	Q5_K_M · 19 GB · 131,072 ctx
6	Nemotron 3.5 Lightning 30B-A3B NVIDIA, hybrid Mamba-2 + MoE, 3B active · launched 11 Aug 2026	LOCAL	0.800	96.7%	51s	IQ4_XS · 18 GB · 131,072 ctx
7	Qwen3 Coder 30B-A3B MoE 30B, 3B active, coding-specialist	LOCAL	0.667	86.0%	144s	Q4_K_M · 18 GB · 65,536 ctx

The trade-off

Quality vs. latency

Up and to the left is better: higher judged quality, lower median time to a finished task. Log-scaled x-axis.



The headline result: the best local models land within **3.3 points of pass rate** and **0.07 of quality** of hosted frontier models, at comparable median latency — and one local model is faster than two of the three frontier endpoints. The gap that remains is concentrated in judged craft, not in whether the task got done.

How the weights were fitted to the card

Every local model is quantized to sit inside 24 GB of VRAM *with its KV cache at the context length it is actually served at* — the quantization is chosen per model rather than applied uniformly, and context is capped where the KV cache would otherwise crowd the weights off the card. Qwen3.6 27B runs Q4_K_M at 98k; Muse Glimmer 30B takes a heavier Q5_K_M at the full 131k; Nemotron 3.5 Lightning, whose hybrid Mamba-2 layers keep its KV cache small, runs IQ4_XS at 131k. Qwen3 Coder is deliberately held to 65k: measured at higher context it needed 22.5 GB and left no headroom on the card. Models are hot-swapped one at a time; only one is resident at any moment.

MODEL	NETDEVOPS	AIOPS	HPC SCHEDULING	INFERENCE SERVING	RAG PIPELINES	TOOL-CALLING LOOPS
GPT-5.6 Sol	100%	100%	100%	100%	100%	100%
Claude Opus 4.8	100%	100%	100%	100%	100%	100%
GPT-5.5	100%	100%	100%	100%	70%	100%
Qwen3.6 27B	100%	100%	50%	100%	100%	100%
Muse Glimmer 30B	100%	100%	50%	100%	100%	100%
Nemotron 3.5 Lightning 30B-A3B	100%	100%	50%	100%	100%	100%
Qwen3 Coder 30B- A3B	75%	100%	100%	100%	90%	100%


Pass rate by domain, pooled across both harnesses. Dashes mark dimensions with no task in that tier for this model.

Same model, two harnesses

Because each model was driven by two different agent harnesses over identical tasks, the harness itself becomes measurable. It matters more than expected: the scaffolding around a model moves pass rate by up to 15 points on the same weights.

Same model, two agent harnesses

Pass rate for each model driven by each harness. Identical tasks, identical model, identical grading.



Disclaimer. This is not an official benchmark result for any model, vendor, product or organisation, and it is not endorsed by any of them. These are the numbers I measured on my own hardware, with my own task battery, my own quantizations and my own serving configuration, on a single day. Sample sizes are small (15 tasks per model per harness), the grading judge is itself a language model, and the local models are quantized — they are not the full-precision weights their publishers released, so these figures understate them by an amount I have not measured. Hosted frontier models run on infrastructure whose configuration I do not control and which can change without notice. Results are not directly comparable to any published leaderboard and should not be read as a general claim about any model's capability. Treat this as one operator's engineering log, not as evidence.

Local models served on a single NVIDIA RTX 3090 (24 GB) over Thunderbolt 4, behind a LiteLLM gateway. Task battery and runner published for reproduction at github.com/mairp/agent-ops-bench.